

SNP calling from RNAseq data

Steve Palumbi Lab

Hopkins Marine Station, Stanford University

Analysis of sequence information from RNA-Seq data. (workshop) SNP-calling and overview of methods to analyze the output. **Steve Palumbi lab at Stanford/Hopkins.**

Part 1: 35 minutes

Overview: from Bam files to genotypes: Steve Palumbi [3 min]

Bowtie2 and SamTools: Bryan Barney and Nathan Truelove [10]

FreeBayes: Noah Rose and Elora López [10]

vcfTools and the 0,1,2 genotype file: Beth Sheets and Megan Morikawa [3]

PCA and FST: Megan Morikawa and Bryan Barney [5]

Part 2: 30 minutes

Mentored file manipulation workshop from fastQ files to 0,1,2 genotype matrix using demo input files

Part 3: 30 minutes


Overview: using genotype data: Steve Palumbi [1 min]

NgsAdmix and linkage: Bryan Barney [7]

Outliers and environmental correlations: Luke Thomas and Nathan Truelove [5]

Somatic mutations: Elora López [5]

dN/dS and eQTLs: Noah Rose [9]



Mapping in Bowtie2


Bowtie2 lines up reads to a reference genome or transcriptome

End-to-End: Uses all the base-pairs

```
Read:      G A C T G G G C G A T C T C G A C T T C G
           | | | | |   | | | | | | | | | | | |
Reference:  G A C T G - - C G A T C T C G A C A T C G
```

Local: Base-pairs at the ends can be discarded

```
Read:      A C G G T T G C G T T A A - T C C G C C A C G
           | | | | | | | | | | | |
Reference:  T A A C T T G C G T T A A A T C C G C C T G G
```



Alignment Score

How similar the read is to the reference

End-to-End Example:

Mismatch = -6

Read Gap = -11

Best alignment score = 0


Local Example:

Mismatch = -6

Read Gap = -11

Base that matches Reference = +2

Best Alignment Score = 2 x Read length



Minimum Alignment Score

- Expressed as a Function of Read Length:

$$f(x) = 0 + -0.6 * x, \text{ where } x \text{ is the read length}$$


- End-to-end alignment mode default is:

$$-0.6 + -0.6 * \text{read length}$$

- For a 50 base-pair read:

$$-0.6 + -30 = -30.6$$

- Default: 5 mismatches/2 read gaps/Combos




Optimize Mapping Parameters

End-to-End Example:

Mismatch = -6

Read Gap = -11

Best alignment score = 0



Optimize Mapping Parameters

--score min

- Changes the default minimum alignment score to be considered valid.

Default: $L, 0, -0.6 = -30.6$ for 50 bp reads

- Optimized for 2 mismatches/1 mismatch and 1 read gap


--score-min $L, 0, -0.36 = -18$



Preset Mapping Parameters

Verify that the preset meets your mapping requirements

- very-fast
- fast
- sensitive
- very-sensitive



Mapping : Bowtie outputs a SAM file

SAM files contain a list of reads, each read will get a series of 'fields' associated with it that describe the mapping result

Col	Field	Type	Regexp/Range	Brief description
1	QNAME	String	[!-?A-~]{1,254}	Query template NAME
2	FLAG	Int	[0,2 ¹⁶ -1]	bitwise FLAG
3	RNAME	String	* [!-()+-<>-~] [!-~]*	Reference sequence NAME
4	POS	Int	[0,2 ³¹ -1]	1-based leftmost mapping POSition
5	MAPQ	Int	[0,2 ⁸ -1]	MAPping Quality
6	CIGAR	String	* ([0-9]+[MIDNSHPX=])+	CIGAR string
7	RNEXT	String	* = [!-()+-<>-~] [!-~]*	Ref. name of the mate/next read
8	PNEXT	Int	[0,2 ³¹ -1]	Position of the mate/next read
9	TLEN	Int	[-2 ³¹ +1,2 ³¹ -1]	observed Template LENgth
10	SEQ	String	* [A-Za-z=.]+	segment SEQuence
11	QUAL	String	[!-~]+	ASCII of Phred-scaled base QUALity+33

Different mapping utilities (bowtie2, BWA, etc) will have different 'additional fields' that you might use for filtering



Mapping : Samtools converts SAM to BAM, sorts, & indexes

- SAM files are human readable plain text
- BAM files are binary versions of SAM that are smaller and easier for the computer to process

Sorting:

- Groups your reads by where in the assembly they mapped


Indexing:

- Makes a table that contains information about:
 - How many reads mapped
 - Where they mapped
 - Reads that didn't map

SNP Calling

711 721 731 741 751 761 771 781 791
GTGGCGAGAAAATGTCGATCGCCATTATGGCCGGCGTGTAGAAAGCGCGTGGTCACAACGTTACCGTTATCGATCCGGTCGAAAAACTGCTGGCAGT
.....G.....
.,,,,,,,,,,,,,,g,,,,,,,,,,,,,c,,,,,,,,,,,,,
.,,.,c.,t.,,,,,,,,,,,,,,g,,,,,,,,,,,,,
.,,.,,t.,,,,,,,,,,,,,,g,,,,,,,,,,,,,
.,,.,,,,,,,,,,,,,,g,,,,,,,,,,,,,
.,,.,,,,,,,,,,,,,,tg,,,,,,,,,,,,,t.,
.,,.,,c.,,,,,,,,,,,,,,g,,,,,,,,,,,,,g,,,,,,,,,,,,,
c.,,.,,,,,,,,,,,,,,g.,a.,,,,,,,,,,,,,,t.,
.,,.,,,,,,,,,,,,,,gc.,,,,,,,,,,,,,,
.,,.,,g.,g.,,,,,,,,,,,,,,
.,,.,,,,,,,,,,,,,,g.,g.,,,,,,,,,,,,,,
.,,.,,,,,,,,,,,,,,g.,g.,,,,,,,,,,,,,,a,c.,,
.,,.,,,,,,,,,,,,,,g.,,,,,,,,,,,,,,g.,,
.,,.,,g.,g.,,,,,,,,,,,,,,c.,,
.,,.,,g.,,,,,,,,,,,,,,
.,,.,,g.,,,,,,,,,,,,,,g.,g.,,,,,,,,,,,,,,t.,,
.,,.,,g.,,,,,,,,,,,,,,
.,,.,,g.,,,,,,,,,,,,,,t.,,,,,,,,,,,,,,t.,,
.,,.,,g.,,,,,,,,,,,,,,c.,,,,,,,,,,,,,,a.,g.,,
.,,.,,g.,,,,,,,,,,,,,,c.,,,,,,,,,,,,,,a.,t.,,
.,c.,,,,,,,,,,,,,,tg.,,,,,,,,,,,,,,g.,g.,,
.,,.,t.,,,,,,,,,,,,,,g.,t.,,,,,,,,,,,,,,a.,,

Haplotype-based approach





Freebayes pipeline (minimal)

- Start with fastq reads
- Map reads (e.g. bowtie2 or hisat2)
- If paired end libraries with PCR amplification, remove PCR duplicates (Picardtools)
- Sort, compress, and index alignments (samtools)
- Call SNPs (Freebayes) Minimal call: `freebayes -f ref.fa *.bam > out.vcf`
- Filter SNPs (vcffilter)



Pros and Cons of Freebayes

- Pros
 - Fast, sophisticated model
 - Easy interface, easy to customize via command line arguments
 - Good support for local multithreading (freebayes-parallel) and cluster parallelization (just split a bed file of your contigs into as many jobs as you like)



Pros and Cons of Freebayes

- Cons
 - Relentlessly haplotype based, so it can sometimes be hard to get just, like, normal biallelic SNPs (this is a feature too)
 - Utilities like vcffilter, vcfallelicprimitives, and vcfbiallelic help
 - Under rapid development, so sometimes tools change or useful features haven't been implemented or documentation is less good



GATK's HaplotypeCaller

- Defines “active regions”
- Determines haplotypes by reassembling the active region
- Determines likelihoods of the haplotypes given the read data
- Assigns sample genotypes
- Outputs VCF or gVCF file




Pros and Cons of GATK

- Pros
 - Extensively documented, lots of support
 - Clear, easy-to-interpret output



Pros and Cons of GATK

- Cons
 - Not as easily customizable, not as easy interface as Freebayes
 - Slower than Freebayes



012 SNP Matrix - format

- Variant call file (.vcf) – list of alleles and their likelihoods
- Use vcftools to convert your filtered SNP file (.vcf) into a 012 matrix
- Each row is a sample, each column is a SNP
 - First column is sample number, starting at 0
- 0 : both copies of reference allele
- 1: heterozygous
- 2: both copies of alternate allele

ex:

0	0	1	0	0
1	0	0	0	2
2	2	0	1	0
3	2	0	0	1

012 SNP Matrix vs. other methods

- 012 genotype calls does not represent uncertainty about genotype
 - Ex: If we only have 2 mapped reads, both the alternate allele at the locus, this could be homozygous alternate or a heterozygote where we did not sample the other allele
- We can remove uncertainty by filtering for SNP calls that we are very confident about
- This is the strategy we are using in the pipeline today

Other methods: using genotype likelihoods

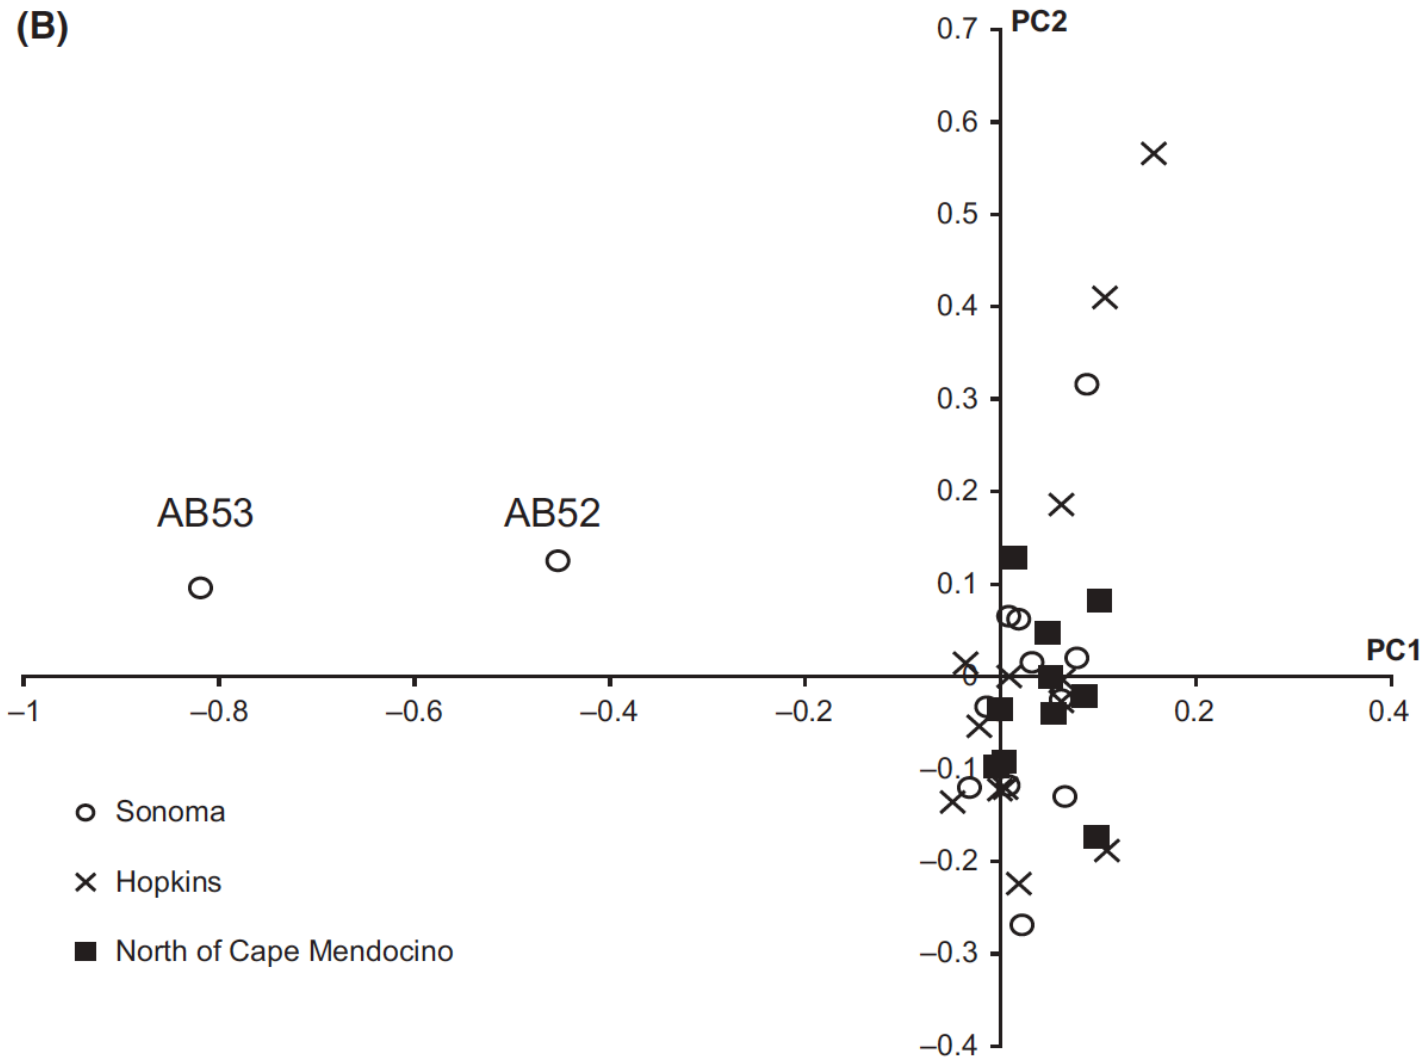
- Fewer programs use this format, examples are GATK++, Angsd

Principle Components Analysis

Taking thousands of SNPs into account simultaneously


- Population structure
 - Detecting outliers
- An axis of variation to compare to environmental variables

Population differentiation




Transcriptome-wide, no differentiation

Clones in natural populations



11 points on PC1 & PC2 from 20 samples

Correlation to environmental variables




PC4 and correlation to environment

F_{ST} analysis

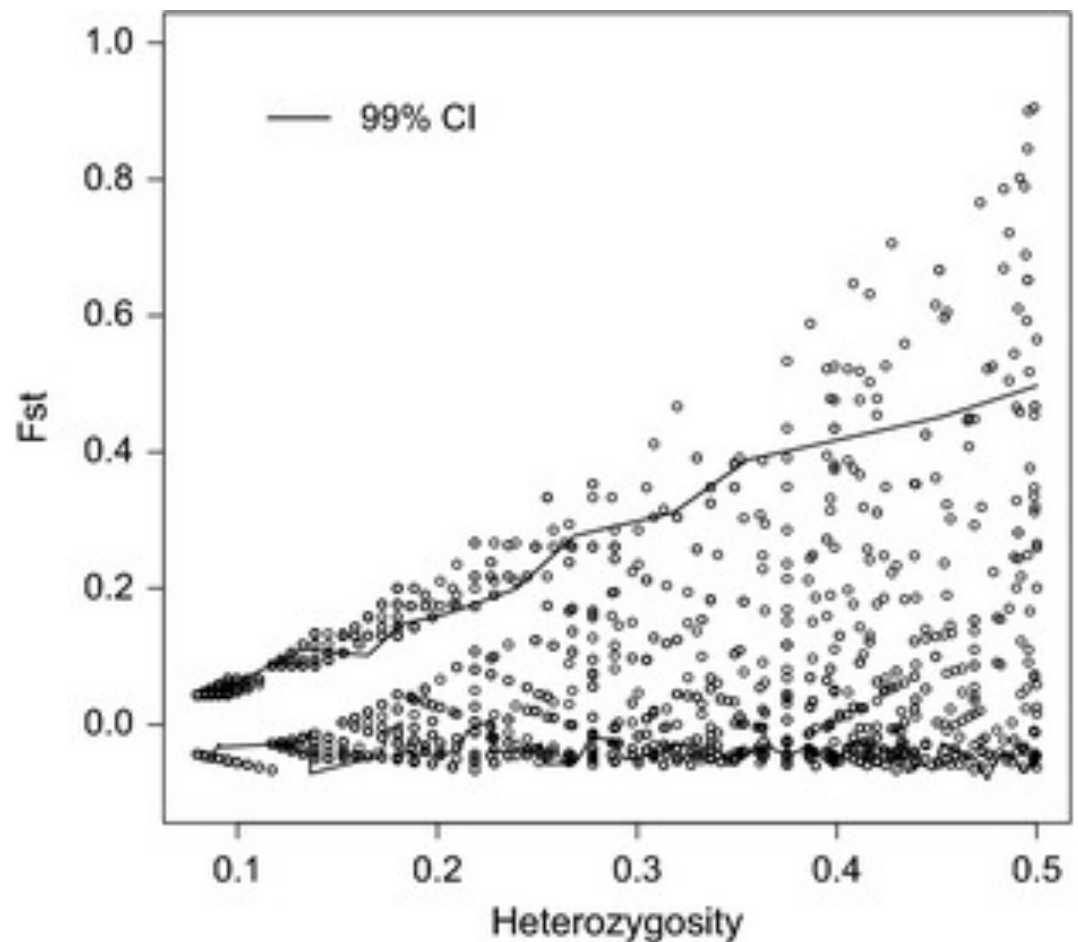
- F_{ST} as an index of genetic differentiation
 - Ranges from 0 (identical) to 1 (completely different)
- Classic measure – can compare to MANY studies
- Wright (1953) F_{ST} vs Weir and Cockerham (1984) F_{ST}
- With transcriptomic levels of data, we need to look at patterns, not necessarily individual loci

F_{ST} analysis

Compare density distributions of pairwise F_{ST} between pops



Compare F_{ST} to heterozygosity to find unusually high F_{ST} loci



Tutorial

https://github.com/bethsheets/SNPcalling_tutorial

Why use transcriptomics?

- Reduced representation
- Focus on parts that matter (protein coding)
- Expression links SNPs to phenotype

Population structuring : NGSadmix

<http://www.popgen.dk/software/index.php/ANGSD>


<http://www.popgen.dk/software/index.php/NgsAdmix>

NGSadmix takes genotype likelihood files and outputs a view of population membership of individuals



Linkage disequilibrium : within contig

- Arises from a lack of recombination between SNPs
- SNPs inherited as a block, not necessarily adjacent



BUT: LD may exist beyond the extent of individual transcripts, across multiple genes on the same chromosome, or across chromosomes!

Linkage disequilibrium : whole genome

Extra data are needed for better understanding of physical linkage at the chromosomal level:

- linkage maps from pedigreed individuals
- OR (*AND* is better!), a well-assembled genome

Pairwise LD calculations for all SNPs throughout assembly

- Generates a matrix of r^2 values for each pair of SNPs


Decision: How to cluster the SNPs, and what cutoff to use?

- single-linkage clustering (A linked to B, B to C, so A,B,C in cluster)
- $r^2 > 0.75$

Linkage disequilibrium : whole genome

“Islands of divergence” or supergenes?

A GO enrichment analysis of linked region may reveal overrepresentation of genes of related function, the classical definition of a supergene





Genome wide scans for selection: identify genomic regions that exhibit signatures of diversifying selection.

- Two main approaches:
 1. **Population differentiation** (PD) approaches
 2. **Ecological association** (EA) approaches

Population differentiation (PD) approaches

- Identifying loci that show unusual allele frequency differentiation among populations




Population differentiation (PD) approaches


- Pros:
 - Can be effective at identifying genes under selection without known phenotypes
 - Does not require *a priori* information concerning the environmental forces that act as selective pressures.
 - Can screen high number of markers to identify candidate genes for further investigation.
- Cons
 - Plagued by false positives and false negatives
 - Candidate outlier loci often vary in pairwise population comparisons, and therefore overall divergence (global F_{ST}) may not detect candidates that are under selection in only a portion of populations.
 - Limited power in detecting balancing selection and various forms of weak divergent selection.

Ecological association (EA) approaches

- Identifying loci with a strong association between allele frequencies and environmental variables



<http://www.sccwrp.org>



<https://sites.google.com/site/nkooyers/>

Ecological association (EA) approaches

- Pros:
 - Uncover selected loci without knowledge of selective environment (can feed the programs lots of data and see what comes out)
 - More powerful than PD approaches
- Cons
 - Requires detailed environmental data
 - If IBD then surprisingly low power (need to account for population structure)
 - High variability among runs

Genome wide scans for selection: Conclusions

- Assist in identifying loci under selection
- Be careful as can lead to numerous false-positives!
- Best used in conjunction with GWAS and linkage-mapping approaches

For reviews on the topic:

- Lotterhos and Whitlock (2015) Mol. Ecol.
- Hoban et al. (2016) Am. Nat.
- Rellstab et al. (2015) Mol. Ecol.
- Narum and Hess (2011) Mol. Ecol. Res.
- Haas and Payseur (2016) Mol. Ecol.

- **Common programs:**

- BayeScan
- Arlequin
- Lositan
- OutFLANK
- PCAdapt
- BAYENV2
- BAYESCENV

Custom Analyses of VCF files

- Identifying somatic variants

ID	REF	ALT	10A	10B	11A	11B	1A	1B	2B	2C	3A	3C	4A	4C	5A	5B	6A	6C	7A	7B	8B	8C	9A	9C
contig105225	T	A	0	1	1	1	0	0	0	0	1	0	1	1	0	1	0	0	1	0	0	0	0	0
contig108861	A	C	0	0	0	1	1	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0
contig114781	A	G	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1
contig138091	T	G	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1
contig142456	C	A	0	0	0	0	0	0	0	1	0	0	0	1	0	0	1	1	0	0	0	0	0	0
contig142456	T	G	0	0	0	0	0	0	0	1	0	0	0	1	0	0	0	0	0	0	0	0	0	0
contig150784	A	C	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
contig151485	A	G	0	0	0	0	1	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0

8 of 214 somatic variants identified in a single coral colony

Green = homozygote

Yellow = heterozygote

How to link SNPs to phenotypes

- Protein coding
- eQTLs

Protein coding change

mRNA



Goal: find protein
coding sequences in
transcripts

Tools:
OrfPredictor
Blastx
GMAP

Imperfect assembly



Protein coding change

mRNA



VCF says:

Contig1 24 T A

Goal: Find amino acid changes due to SNP variants

A..S..M..W..G..T..Y..F..S..W..T..

ORF

Tools:

Biopython

Snpeff

TTT

TTA

A..S..M..W..G..T..Y..**F**..S..W..T..


A..S..M..W..G..T..Y..**L**..S..W..T..

Questions to ask with protein coding SNPs

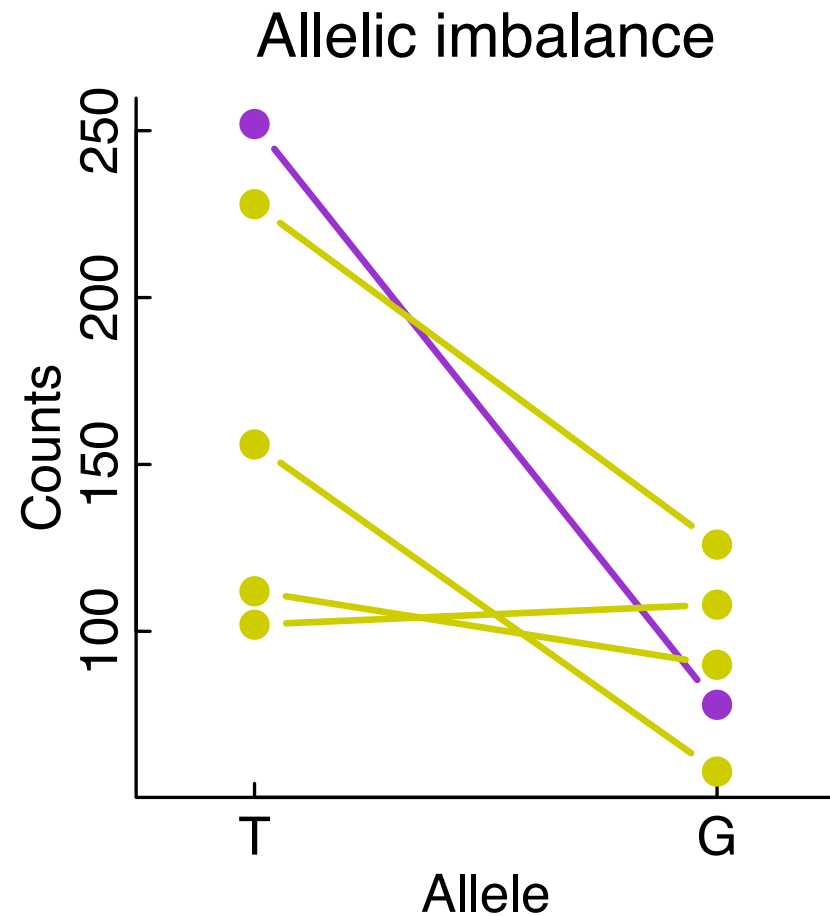

- Does a protein sequence show strong differences between populations or species?
- Do non-synonymous variants fall in residues known to be important from multi-alignment conservation or crystallography?
- Are proteins evolving to be different more quickly than you would expect from genetic drift?
 - McDonald Kreitman test for multiple samples from two species
 - dN/dS (rate of non-synonymous change over rate of synonymous change) is a general index of the rate of protein evolution

How to detect eQTLs:

SNPs that are correlated with expression



eQTLs explain variation within and between species



Questions to ask with eQTL SNPs (eSNPs)

- Do interesting genes show strong differences in expression between populations or species?
- Are these genes involved in gene networks or pathways that you are interested in?
- Do many genes in the same pathway show expression changes in the same direction?